

SoME Testing

v2 (SoME MVP v2)

- Core: d_model=256, num_layers=4, num_heads=4
- MoE: num_experts=64, top_k=4, d_ffn=1024, alpha=0.01, beta=0.001, delta=0.001, theta=200
- Data: seq_len=256, batch_size=64, vocab_size=8192
- Train: epochs=3, lr=8e-4, AdamW (wd=0.1), AMP + torch.compile
- Size: Total ~72.80M, Trainable ~5.52M
- Result (last epoch): val loss 0.3454, ppl 1.41

v3 (SoME MVP v3)

- Core: d_model=384, num_layers=8, num_heads=6
- MoE: num_experts=32, top_k=4, d_ffn=1024, alpha=0.01, beta=0.001, delta=0.001, theta=200
- Data: seq_len=256, batch_size=64, vocab_size=8192
- Train: epochs=2, lr=8e-4, AdamW (wd=0.1), AMP + torch.compile
- Size: Total ~213.91M, Trainable ~12.23M
- Result (epoch 2): val loss 0.1017, ppl 1.11

v4 (SoME MVP v4)

- Core: d_model=384, num_layers=8, num_heads=6
- MoE: num_experts=64, top_k=2, d_ffn=1536, alpha=0.01, beta=0.001, delta=0.001, theta=200
- Data: seq_len=256, batch_size=512, vocab_size=8192
- Train: epochs=1, lr=8e-4, AdamW (wd=0.1), AMP + torch.compile
- Size: Total ~617.19M, Trainable ~12.23M
- Result (epoch 1): Val Loss = 1.9435, ppl 6.98

V5

- Core: D_MODEL = 384, NUM_HEADS= 6, NUM_LAYERS = 8
- MoE: num_experts: 32, d_ffn: 1024, top_k: 4, theta: 200
- Data: SEQ_LEN = 256, BATCH_SIZE = 64, VOCAB_SIZE = 8192
- Train: Epoch 2: lr=8e-4, AdamW (wd=0.1), AMP + torch.compile
- Size: Total parameters: 213.91M, Trainable parameters: 12.23M
- Results: (Epoch 2) Train Loss = 0.2041, Val Loss = 0.1017, Val Perplexity = 1.11

V6 (upcoming)

- Core: d_model=384, num_layers=10, num_heads=8
- MoE: num_experts=128, top_k=4, d_ffn=1536, alpha=0.01, beta=0.001, delta=0.001, theta=200
- Data: seq_len=256, batch_size=256, vocab_size=8192
- Train: epochs=3, lr=8e-4, AdamW (wd=0.1), AMP + torch.compile

- Size: Total parameters: 1526.11M, Trainable parameters: 13.71M
- Results: Train Loss = 2.3146, Val Loss = 2.2723, Val Perplexity = 9.70

What's Next:

1. Use the Full Dataset: Train the model on the entire TinyStories dataset for a more extended period.
2. Analyze Expert Specialization and Dynamics: The core hypothesis of SoME is that experts will self-organize into "knowledge galaxies." Now is the time to verify this.
 - a. Visualize the Key Store: After training, extract the final key_store tensor from one of the SOMELayers. Use dimensionality reduction techniques like t-SNE or PCA to plot the 32 expert keys in a 2D space. Are there distinct clusters?
 - b. Track Key Movement: Log the positions of the keys at different stages of training (e.g., after epoch 1 and epoch 2). Visualizing this can show how the "knowledge galaxies" form over time.
 - c. Analyze Expert Usage: Plot a histogram of the usage_count buffer.
 - i. Is the load balanced? A relatively even distribution is a good sign.
 - ii. Are there "generalist" experts? A few experts with very high usage counts could be acting as stable "galactic centers," as your theory predicts.
 - iii. Are any experts being pruned? Check if any experts consistently fall below the theta threshold and are being decayed.

Ablation Studies

A. Fix backbone (use v3 core), vary MoE

- v3 core + 32e, k=4 (current) → baseline
- v3 core + 64e, k=4 → tests if v3 still wins when experts increase (watch for renewed imbalance).
- v3 core + 32e, k=8 → holds experts constant and increases mixture breadth per token; often improves ppl if capacity allows.
- Readouts: per-layer gate entropy, load Gini, % dropped tokens, validation ppl.

B. Fix experts (use 32e, k=4), vary backbone

- Shrink to v2 core (4×256) to confirm the bulk of the improvement was core capacity.
- Middle ground (6×320) to find a sweet spot for compute vs. ppl.
- Readouts: ppl vs. total params curve; plot loss vs. trained tokens to see sample-efficiency.

C. Isolate SOME's in-inference adaptation

Run two modes on the same trained checkpoint:

- Frozen keys (disable Query/Peer/Repulsive updates) vs. SOME active (current $\alpha/\beta/\theta$).

- Evaluate under a short domain shift (TinyStories → out-of-domain subset) and report ppl over time and post-shift validation. This directly shows the benefit of address-only plasticity.

D. Stress-test routing stability

- Key norm control: L2-normalize keys after each update; add a learnable temperature τ in routing softmax and target a gate-entropy band (e.g., 1.5–2.5 nats).
- Clip Δk (e.g., global norm or per-step cap) and consider a small EMA over keys to smooth updates.
- Usage-aware α floor: don't let α decay below ~0.005 so cold experts can still adapt.

What each knob does

1) Backbone capacity (dense transformer): These change the shared capacity (applies to all tokens).

1. `d_model` (hidden size)
 - Quality: $\uparrow\uparrow$ (usually the strongest single dense knob).
 - Compute per token: \uparrow (attention, MLP scale with $\sim d^2$).
 - Memory: \uparrow (activations, params/opt states).
 - Range: 320–448 on A100-80G for fast runs.
 - Tip: Increase d before layers if your model is shallow (<8).
2. `num_layers` (depth)
 - Quality: $\uparrow\uparrow$ (great bang-per-FLOP at this scale).
 - Compute per token: \uparrow linearly with layers.
 - Memory: \uparrow (activations).
 - Range: 6–10 is a good sweet spot for sub-hour runs.
3. `num_heads`
 - Quality: \uparrow (helps mixing, but smaller effect than $d/layers$).
 - Compute/Memory: mild \uparrow (within same `d_model`).
 - Range: keep divisors of `d_model`; 6–8 for $d=384–448$.

2) MoE knobs (capacity vs compute): These control sparse capacity and routing compute.

1. `num_experts`
 - Quality: \uparrow (more niches); diminishing returns if router underpowered.
 - Compute per token: ~neutral if `top_k` fixed; memory \uparrow for expert weights (but frozen experts don't add optimizer state).
 - Range: 24–64. For speed, 32 is a sweet spot.
2. `top_k`
 - Quality: \uparrow with larger k (more expert collaboration).
 - Compute per token: \uparrow linearly (k extra FFNs).
 - Memory: small \uparrow .
 - Rule of thumb: If you halve k (4→2), increase `d_ffn` a bit (e.g., 1024→1536) or add LoRA to recover quality at much lower cost.
3. `d_ffn` (expert FFN width)

- Quality: \uparrow with larger FFN.
 - Compute per token: \uparrow (MLP dominates).
 - Tandem with top_k: tune as a product: effective capacity $\approx k \times d_{ffn}$.
4. Experts trainability (frozen vs LoRA vs full)
- Frozen: fastest, smallest optimizer state; rely on routing to specialize.
 - LoRA on experts (e.g., $r=4-8$): big ROI on quality, tiny overhead.
 - Full train: best quality but heavy optimizer/memory/time.

3) Router & self-organization (your special sauce): Affects specialization, stability, and some overhead.

1. Query network (depth/width; e.g., 1-layer linear vs 2-layer MLP $d \rightarrow h \rightarrow d$)
 - Quality: \uparrow with 2-layer (GELU), esp. with more experts.
 - Compute: slight \uparrow . Worth it.
2. Self-org hyperparams: alpha (attraction), beta (repulsion), delta (decay), theta (low-usage threshold)
 - Quality/Stability:
 - i. Higher β early spreads experts (prevents collapse).
 - ii. Moderate α consolidates niches.
 - iii. δ/θ push underused experts to explore.
 - Best practice: β -anneal (e.g., $3e-3 \rightarrow 5e-4$ over training).
3. Key update cadence
 - Every step vs every other step / token subsampling (e.g., 50%):
 - i. Speed: improves 10–20% with minimal quality hit.
 - ii. Use subsampling to buy throughput if time-bound.
4. Eval behavior
 - Never update keys during eval. (Keeps validation stationary and honest.)
However the system is supposed to adapt in real-time, so hypothetically it is fair to allow update keys during eval.
5. Top-k warmup
 - Start with $k=1$ for 10–20% steps \rightarrow ramp to final k .
 - Stability: reduces early routing chaos; often improves final loss.
6. Usage/balance logging
 - Track per-layer entropy and Gini of expert usage.
 - Symptoms: low entropy + high Gini \rightarrow dominance/collapse; increase β or extend $k=1$ warmup.

4) Sequence length, batch, and accumulation: The main speed vs memory levers.

1. seq_len (train)
 - Quality: mild \uparrow if task needs long context.
 - Compute: $O(L^2)$ for attention — expensive.
 - Strategy: Keep L modest (320–384) to go much faster; extrapolate at inference with RoPE/ALiBi.
2. Micro-batch size

- Speed: ↑↑—this is how you use your 80GB.
 - Memory: ↑ (activations).
 - Target: raise until ~60–70 GB used; then set grad_accum to hit your global_batch_tokens target (e.g., 120k–180k tokens/step).
3. Global batch tokens (micro_batch × seq_len × grad_accum)
 - Throughput: scales well until you saturate data pipeline.
 - Stability: don't overshoot LR/BS scaling; keep LR in check.

5) Optimization & schedule:

1. LR (peak) and schedule (warmup + cosine)
 - Quality/Stability: warmup 3–5% of steps is usually sweet; cosine is safe.
 - Note: If you surge batch a lot, LR may need a slight downscale.
2. Weight decay
 - 0.05–0.1 typical for small LM/TinyStories.
3. Aux balance losses
 - I'm using self-org ($\alpha/\beta/\delta/\theta$). If you add a classic MoE load-balancing loss, use a tiny coeff to avoid fighting the self-org mechanism.

6) System throughput knobs (make the GPU sweat):

1. Precision: BF16 autocast.
2. FlashAttention / SDPA: must be on.
3. Fused AdamW (fused=True if available) or Apex FusedAdam.
4. torch.compile(mode="reduce-overhead") after ~50 warmup steps.
5. Disable activation checkpointing unless you truly need longer seq_len.
6. Data loader: pre-tokenize; num_workers=8–16, pin_memory=True, prefetch_factor=2–4, non_blocking=True.
7. Optional: CUDA Graphs if shapes are static (removes launch overhead).

7) Data & tokenizer:

1. Tokenizer/vocab (e.g., 8k)
 - Quality: can help; smaller vocab reduces softmax cost.
2. Packing/bucketing
 - Speed: ↑ by cutting padding; more real tokens per step.

8) Interactions that matter (rules of thumb):

1. top_k × d_ffn: keep their product roughly constant to hold MoE compute steady.
 - Example: k=4, FFN=1024 ≈ k=2, FFN=1536 (the latter is faster).
2. seq_len × micro_batch × grad_accum: fit VRAM via micro-batch first; avoid growing L unless needed.
3. Experts trainability: Frozen + LoRA often matches most of the gain of full finetune at a fraction of cost.
4. Router depth: A small 2-layer MLP router often pays for itself at E≥32.