

SoME v8.1 Full Grafting

```
# =====
# SoME v8.1: The Full Graft (Complete)
# =====

# --- Cell 1: Dependencies ---
import subprocess
import sys

def install_deps():
    print("Installing dependencies...")
    subprocess.check_call([sys.executable, "-m", "pip", "install", "-q", "transformers", "datasets",
    "tokenizers", "accelerate", "tic-toc"])

try:
    import transformers
except ImportError:
    install_deps()

import torch
import torch.nn as nn
import torch.nn.functional as F
from torch.utils.data import DataLoader, Dataset
from transformers import AutoModelForCausalLM, AutoTokenizer
from datasets import load_dataset
from tqdm import tqdm
import math
import os
import numpy as np
import gc
import random

# Setup Compute
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
print(f"Using device: {device}")

# Enable Ampere Optimizations
if torch.cuda.is_available() and torch.cuda.get_device_capability(0)[0] >= 8:
    print("Ampere+ GPU detected. Enabling TF32.")
    torch.set_float32_matmul_precision('high')
    torch.backends.cudnn.benchmark = True

# =====
```

```

# --- Cell 2: The Harvester (Full Graft) ---
# =====

def harvest_qwen_full(model_name="Qwen/Qwen2.5-0.5B", slice_size=128):
    """
    Downloads Qwen, extracts Embeddings and FFNs, and prepares them for SoME.
    Returns: expert_library, embeddings_tensor, config
    """

    print(f"\n📝 Harvesting COMPLETE organs from {model_name}...")

    # Load source model on CPU
    hf_model = AutoModelForCausalLM.from_pretrained(
        model_name,
        torch_dtype=torch.float16,
        device_map="cpu",
        trust_remote_code=True
    )

    # 1. Harvest Embeddings
    # We clone them to detach from the HF model graph
    print(" -> Extracting Embeddings...")
    embeddings = hf_model.model.embed_tokens.weight.data.clone() # [Vocab, Dim]
    vocab_size, d_model = embeddings.shape

    # 2. Harvest Experts (FFNs)
    print(" -> Slicing FFNs (Target Slice: {slice_size})...")
    expert_library = []
    layers = hf_model.model.layers

    for i, layer in enumerate(tqdm(layers, desc="Slicing Layers")):
        mlp = layer.mlp
        # Qwen weights
        w_gate = mlp.gate_proj.weight.data
        w_up = mlp.up_proj.weight.data
        w_down = mlp.down_proj.weight.data

        intermediate_size = w_gate.shape[0]
        num_slices = intermediate_size // slice_size

        for j in range(num_slices):
            start = j * slice_size
            end = start + slice_size

            # Slice SwiGLU components

```

```

# We convert to FP32 for safer training dynamics, though BF16 is also fine
expert_library.append({
    'w_gate': w_gate[start:end, :].clone(),
    'w_up': w_up[start:end, :].clone(),
    'w_down': w_down[:, start:end].clone(),
    'source_layer': i
})

print(f"✅ Harvest Complete. Experts: {len(expert_library)}. Embeddings: {embeddings.shape}")

# Cleanup
del hf_model
gc.collect()
torch.cuda.empty_cache()

return expert_library, embeddings, d_model

# =====
# --- Cell 3: Data Pipeline (Qwen Tokenizer) ---
# =====

class LanguageModelDataset(Dataset):
    def __init__(self, tokenized_data, pad_token_id: int, eos_token_id: int = None):
        self.data = tokenized_data
        self.pad_token_id = pad_token_id
        self.eos_token_id = eos_token_id

    def __len__(self):
        return len(self.data)

    def __getitem__(self, idx):
        item = self.data[idx]
        input_ids = torch.tensor(item["input_ids"], dtype=torch.long)

        if "attention_mask" in item:
            attention_mask = torch.tensor(item["attention_mask"], dtype=torch.long)
        else:
            attention_mask = (input_ids != self.pad_token_id).long()

        targets = input_ids.clone()
        targets[:-1] = input_ids[1:]
        targets[-1] = -100 # Ignore last token prediction

```

```

if self.pad_token_id is not None:
    targets[targets == self.pad_token_id] = -100

return input_ids, targets, attention_mask

def prepare_data_qwen(config, model_name="Qwen/Qwen2.5-0.5B"):
    print("\n--- Data Preparation (Qwen Tokenizer) ---")

    # 1. Load Pre-trained Tokenizer
    print(f"Loading tokenizer from {model_name}...")
    tokenizer = AutoTokenizer.from_pretrained(model_name, trust_remote_code=True)
    if tokenizer.pad_token is None:
        tokenizer.pad_token = tokenizer.eos_token

    print(f"Tokenizer Vocab Size: {tokenizer.vocab_size}")

    # 2. Load Dataset
    print("Loading TinyStories...")
    full_dataset = load_dataset("roneneldan/TinyStories", streaming=False)
    train_subset = full_dataset['train'].select(range(config['data']['train_subset_size']))
    val_subset = full_dataset['validation'].select(range(config['data']['val_subset_size']))

    # 3. Tokenize
    def tokenize_function(examples):
        # We assume the tokenizer handles EOS automatically or we append it
        return tokenizer(examples["text"], truncation=True, padding="max_length",
                         max_length=config['model']['seq_len'], return_tensors="pt")

    print("Tokenizing...")
    tokenized_train = train_subset.map(tokenize_function, batched=True,
                                       remove_columns=["text"], num_proc=os.cpu_count())
    tokenized_val = val_subset.map(tokenize_function, batched=True,
                                   remove_columns=["text"], num_proc=os.cpu_count())

    train_dataset = LanguageModelDataset(tokenized_train,
                                         pad_token_id=tokenizer.pad_token_id)
    val_dataset = LanguageModelDataset(tokenized_val, pad_token_id=tokenizer.pad_token_id)

    train_loader = DataLoader(train_dataset, batch_size=config['data']['batch_size'],
                             shuffle=True, drop_last=True, num_workers=2, pin_memory=True)
    val_loader = DataLoader(val_dataset, batch_size=config['data']['batch_size'],
                           drop_last=False, num_workers=2, pin_memory=True)

return train_loader, val_loader, tokenizer

```

```

def calculate_gini(usage_counts):
    counts = usage_counts.cpu().to(torch.float32).numpy()
    if np.sum(counts) == 0: return 0.0
    counts = np.sort(counts)
    n = len(counts)
    index = np.arange(1, n + 1)
    return (np.sum((2 * index - n - 1) * counts)) / (n * np.sum(counts))

# =====
# --- Cell 4: SoME v8.1 Architecture ---
# =====

class GraftedExpert(nn.Module):
    """Frozen SwiGLU Expert initialized from Qwen weights."""
    def __init__(self, d_model, d_ffn, weights):
        super().__init__()
        self.w_gate = nn.Linear(d_model, d_ffn, bias=False)
        self.w_up = nn.Linear(d_model, d_ffn, bias=False)
        self.w_down = nn.Linear(d_ffn, d_model, bias=False)

        # Initialize
        self.w_gate.weight.data = weights['w_gate'].to(torch.float32)
        self.w_up.weight.data = weights['w_up'].to(torch.float32)
        self.w_down.weight.data = weights['w_down'].to(torch.float32)

        # Freeze
        for param in self.parameters():
            param.requires_grad = False

    def forward(self, x):
        # Standard SwiGLU: Down(Swish(Gate) * Up)
        return self.w_down(F.silu(self.w_gate(x)) * self.w_up(x))

class SOMELayer(nn.Module):
    def __init__(self, d_model, some_config, expert_library_subset):
        super().__init__()
        self.d_model = d_model
        self.numExperts = len(expert_library_subset)
        self.top_k = some_config['top_k']

        # Heuristics
        self.alpha = some_config['alpha']
        self.beta = some_config['beta']

```

```

self.delta = some_config['delta']
self.respawn_threshold = some_config.get('respawn_threshold', 0.1)
self.theta_percentile = some_config['theta_percentile']
self.warmup_steps = some_config['warmup_steps']
self.ema_decay = some_config['ema_decay']
self.ablation_flags = some_config.get('ablation_flags', {'use_alpha': True, 'use_beta': True,
'use_delta': True})

# Router (Trainable MLP)
mult = float(some_config.get("router_mlp_mult", 2.0))
hidden = int(d_model * mult)
self.query_network = nn.Sequential(
    nn.Linear(d_model, hidden),
    nn.GELU(),
    nn.Linear(hidden, d_model)
)

# Init Router
for m in self.query_network.modules():
    if isinstance(m, nn.Linear):
        nn.init.normal_(m.weight, mean=0.0, std=0.02)
        if m.bias is not None: nn.init.zeros_(m.bias)

# Key Store (Heuristic)
keys = torch.randn(self.num_experts, d_model)
self.register_buffer("key_store", F.normalize(keys, p=2, dim=-1))

# Stats
self.register_buffer("usage_count", torch.zeros(self.num_experts))
self.register_buffer("usage_mass", torch.zeros(self.num_experts))
self.register_buffer("steps", torch.zeros(), dtype=torch.long)
self.register_buffer("dead_expert_count", torch.zeros(), dtype=torch.long)

# Experts (Frozen)
self.experts = nn.ModuleList()
for weights in expert_library_subset:
    d_ffn = weights['w_gate'].shape[0]
    self.experts.append(GraftedExpert(d_model, d_ffn, weights))

if self.top_k > 1:
    self.register_buffer("peer_pull_indices", torch.combinations(torch.arange(self.top_k),
r=2))

def forward(self, x, temperature=1.0):

```

```

batch_size, seq_len, _ = x.shape
x_flat = x.view(-1, self.d_model)

with torch.cuda.amp.autocast(enabled=False):
    # 1. Query
    queries = F.normalize(self.query_network(x_flat.float()), p=2, dim=-1)

    # 2. Similarity
    scores = torch.matmul(queries, self.key_store.t())
    top_k_scores, top_k_indices = torch.topk(scores, self.top_k, dim=-1)

    # 3. Gating
    gating_weights = F.softmax(top_k_scores / float(temperature), dim=-1)

    # 4. Dispatch
    flat_indices = top_k_indices.view(-1)
    sorted_indices, perm_map = torch.sort(flat_indices)
    unique_ids, counts = torch.unique_consecutive(sorted_indices, return_counts=True)

    flat_x = x_flat.repeat_interleave(self.top_k, dim=0)
    perm_x = flat_x[perm_map]
    split_x = torch.split(perm_x, counts.tolist())

    # 5. Compute
    outputs = []
    for i, expert_id in enumerate(unique_ids):
        outputs.append(self.experts[expert_id](split_x[i]))

    cat_out = torch.cat(outputs, dim=0)

    # 6. Reassemble
    inv_perm = torch.argsort(perm_map)
    ordered_out = cat_out[inv_perm]

    # 7. Weight
    gating_weights_f = gating_weights.to(ordered_out.dtype)
    weighted = (ordered_out.view(-1, self.top_k, self.d_model) *
    gating_weights_f.unsqueeze(-1)).sum(1)
    final_output = weighted.view(batch_size, seq_len, self.d_model)

    return x + final_output, queries, top_k_indices, gating_weights

@torch.no_grad()
def update_keys(self, queries, top_k_indices, gating_weights):

```

```

self.steps += 1
num_tokens = top_k_indices.shape[0]
if num_tokens == 0: return

# Update Usage
flat_idx = top_k_indices.reshape(-1)
unique_indices, counts = torch.unique(flat_idx, return_counts=True)
counts_f = counts.to(torch.float32)

self.usage_count.mul_(self.ema_decay)
self.usage_count.index_add_(0, unique_indices, (1.0 - self.ema_decay) * counts_f)

inertia_source = self.usage_count

# Attraction & Peer Pull
if self.ablation_flags.get('use_alpha', True):
    for i in range(self.top_k):
        indices = top_k_indices[:, i]
        inertia = 1.0 + inertia_source[indices]
        alpha_effective = self.alpha / inertia.unsqueeze(-1)
        update_vec = queries - self.key_store[indices]
        self.key_store.index_add_(0, indices, alpha_effective * update_vec)

if self.top_k > 1 and self.ablation_flags.get('use_beta', True):
    indices_i = top_k_indices[:, self.peer_pull_indices[:, 0]].reshape(-1)
    indices_j = top_k_indices[:, self.peer_pull_indices[:, 1]].reshape(-1)
    keys_i, keys_j = self.key_store[indices_i], self.key_store[indices_j]

    inertia_i = (1.0 + inertia_source[indices_i]).unsqueeze(-1)
    inertia_j = (1.0 + inertia_source[indices_j]).unsqueeze(-1)
    beta_effective = self.beta / torch.max(inertia_i, inertia_j)

    update_vec_i = beta_effective * (keys_j - keys_i)
    update_vec_j = beta_effective * (keys_i - keys_j)

    self.key_store.index_add_(0, indices_i, update_vec_i)
    self.key_store.index_add_(0, indices_j, update_vec_j)

# Clamp Fix
current_norms = self.key_store.norm(p=2, dim=-1, keepdim=True)
clamp_mask = (current_norms > 1.0)
if clamp_mask.any():
    self.key_store = torch.where(clamp_mask, self.key_store / current_norms,
                                self.key_store)

```

```

# Decay
if self.steps > self.warmup_steps and self.ablation_flags.get('use_delta', True):
    active_usage_counts = self.usage_count[self.usage_count > 0]
    if active_usage_counts.numel() > 0:
        dynamic_theta = torch.quantile(active_usage_counts.float(), self.theta_percentile)
        low_usage_mask = self.usage_count < dynamic_theta
        self.key_store[low_usage_mask] *= (1.0 - self.delta)

# Phoenix
key_norms = self.key_store.norm(p=2, dim=-1)
dead_mask = key_norms < self.respawn_threshold
num_dead = dead_mask.sum().item()
self.dead_expert_count.fill_(num_dead)

if num_dead > 0:
    flat_queries = queries.view(-1, self.d_model)
    if flat_queries.size(0) >= num_dead:
        rand_indices = torch.randperm(flat_queries.size(0),
device=queries.device)[:num_dead]
        new_keys = flat_queries[rand_indices].clone()
        new_keys = F.normalize(new_keys + torch.randn_like(new_keys) * 0.01, p=2, dim=-1)
        self.key_store[dead_mask] = new_keys
        self.usage_count[dead_mask] = 0.0

class SOMETransformer(nn.Module):
    def __init__(self, model_config, some_config, expert_library, pretrained_embeddings=None):
        super().__init__()
        self.d_model = model_config['d_model']
        self.num_layers = model_config['num_layers']

    # 1. GRAFTED EMBEDDINGS
    if pretrained_embeddings is not None:
        vocab_size, dim = pretrained_embeddings.shape
        self.embedding = nn.Embedding(vocab_size, dim)
        # Initialize with Qwen embeddings
        self.embedding.weight.data = pretrained_embeddings.to(torch.float32)
        # We allow embeddings to fine-tune to TinyStories distribution
    else:
        self.embedding = nn.Embedding(model_config['vocab_size'], self.d_model)

    self.pos_encoder = nn.Embedding(model_config['seq_len'], self.d_model)

    self.layers = nn.ModuleList()

```

```

# Distribute experts
random.shuffle(expert_library)
totalExperts = len(expert_library)
experts_per_layer = totalExperts // self.num_layers

for i in range(self.num_layers):
    start = i * experts_per_layer
    end = start + experts_per_layer
    subset = expert_library[start:end]

    block = nn.ModuleDict({
        'attn': nn.MultiheadAttention(self.d_model, model_config['num_heads'],
batch_first=True),
        'norm1': nn.LayerNorm(self.d_model),
        'norm2': nn.LayerNorm(self.d_model),
        'some': SOMELayer(self.d_model, some_config, subset)
    })
    self.layers.append(block)

self.fc_out = nn.Linear(self.d_model, model_config['vocab_size'])
self.norm_f = nn.LayerNorm(self.d_model)

# Tie Output Head weights to Embedding weights?
# Standard practice for Qwen/Llama. Let's do it for better convergence.
self.fc_out.weight = self.embedding.weight

def forward(self, x, attention_mask=None, temperature=1.0):
    b, seq = x.shape
    positions = torch.arange(seq, device=x.device).unsqueeze(0)
    x = self.embedding(x) + self.pos_encoder(positions)

    all_queries, all_indices, all_gates = [], [], []
    causal_mask = torch.triu(torch.ones(seq, seq, device=x.device), diagonal=1).bool()

    if attention_mask is not None:
        key_padding_mask = (attention_mask == 0)
    else:
        key_padding_mask = None

    for layer in self.layers:
        # Attn
        norm_x = layer['norm1'](x)
        attn_out, _ = layer['attn'](norm_x, norm_x, norm_x,

```

```

        attn_mask=causal_mask,
        key_padding_mask=key_padding_mask,
        is_causal=True)
x = x + attn_out

# SoME (Output is x + residual)
norm_x = layer['norm2'](x)
some_out, q, idx, g = layer['some'](norm_x, temperature=temperature)

# Reconstruct correct residual flow
expert_residual = some_out - norm_x
x = x + expert_residual

all_queries.append(q)
all_indices.append(idx)
all_gates.append(g)

return self.fc_out(self.norm_f(x)), all_queries, all_indices, all_gates

@torch.no_grad()
def update_all_keys(self, all_queries, all_indices, all_gates, token_mask=None):
    if token_mask is not None:
        if token_mask.dim() == 2: token_mask = token_mask.reshape(-1)
        token_mask = token_mask.to(dtype=torch.bool, device=all_indices[0].device)

    for layer, q, idx, g in zip(self.layers, all_queries, all_indices, all_gates):
        if token_mask is not None:
            if q is not None: q = q[token_mask]
            if idx is not None: idx = idx[token_mask]
            if g is not None: g = g[token_mask]

        layer['some'].update_keys(q, idx, g)

# =====
# --- Cell 5: Execution ---
# =====

def train_epoch(model, dataloader, optimizer, criterion, scheduler, current_temp, vocab_size):
    model.train()
    total_loss = 0
    scaler = torch.cuda.amp.GradScaler()
    progress_bar = tqdm(dataloader, desc="Training", leave=False)

    for input_ids, targets, attention_mask in progress_bar:

```

```

    input_ids, targets, attention_mask = input_ids.to(device, non_blocking=True),
    targets.to(device, non_blocking=True), attention_mask.to(device, non_blocking=True)

    with torch.cuda.amp.autocast():
        logits, queries, indices, gates = model(input_ids, attention_mask=attention_mask,
temperature=current_temp)
        loss = criterion(logits.view(-1, vocab_size), targets.view(-1))

        optimizer.zero_grad(set_to_none=True)
        scaler.scale(loss).backward()
        scaler.unscale_(optimizer)
        torch.nn.utils.clip_grad_norm_(model.parameters(), 1.0)
        scaler.step(optimizer)
        scaler.update()
        scheduler.step()

# Heuristics
model_state = model._orig_mod if hasattr(model, "_orig_mod") else model
valid_mask = (targets.view(-1) != -100)
model_state.update_all_keys(queries, indices, gates, token_mask=valid_mask)

total_loss += loss.item()
progress_bar.set_postfix({'loss': f'{loss.item():.4f}'})
return total_loss / len(dataloader)

def evaluate_epoch(model, dataloader, criterion, vocab_size, eval_temp):
    model.eval()
    total_loss = 0
    with torch.no_grad():
        for input_ids, targets, attention_mask in tqdm(dataloader, desc="Eval", leave=False):
            input_ids, targets, attention_mask = input_ids.to(device), targets.to(device),
attention_mask.to(device)
            with torch.cuda.amp.autocast():
                logits, _, _, _ = model(input_ids, attention_mask=attention_mask,
temperature=eval_temp)
                loss = criterion(logits.view(-1, vocab_size), targets.view(-1))
                total_loss += loss.item()
    return total_loss / len(dataloader)

def main():
    # 1. Harvest from Qwen
    expert_library, embeddings, qwen_d_model = harvest_qwen_full("Qwen/Qwen2.5-0.5B",
slice_size=128)
    qwen_vocab_size = embeddings.shape[0]

```

```

# 2. Config
config = {
    "run_name": "v8.1_FullGraft_Qwen0.5B",
    "data": { "train_subset_size": 20000, "val_subset_size": 2000, "batch_size": 16 },
    "model": {
        "vocab_size": qwen_vocab_size,
        "d_model": qwen_d_model,
        "num_heads": 14,
        "num_layers": 6,
        "seq_len": 512
    },
    "some_layer": {
        "top_k": 4,
        "alpha": 0.015,
        "beta": 0.001,
        "delta": 0.005,
        "respawn_threshold": 0.1,
        "theta_percentile": 0.05,
        "warmup_steps": 200,
        "ema_decay": 0.995,
        "router_mlp_mult": 2.0,
        "ablation_flags": {"use_alpha": True, "use_beta": True, "use_delta": True}
    },
    "training": { "num_epochs": 2, "learning_rate": 3e-4, "training_temp": 1.0, "eval_temp": 1.0 }
}

```

3. Data (Using Qwen Tokenizer)

```
train_loader, val_loader, tokenizer = prepare_data_qwen(config, "Qwen/Qwen2.5-0.5B")
```

4. Init Model

```

print("Initializing Full Graft SoME...")
model = SOMETransformer(config['model'], config['some_layer'], expert_library,
pretrained_embeddings=embeddings).to(device)
model = torch.compile(model)

```

5. Stats

```

total_params = sum(p.numel() for p in model.parameters())
trainable_params = sum(p.numel() for p in model.parameters() if p.requires_grad)
print(f"Total Params: {total_params/1e6:.2f}M")
print(f"Trainable Params: {trainable_params/1e6:.2f}M
({trainable_params/total_params*100:.2f}%)"

```

6. Train

```

optimizer = torch.optim.AdamW([p for p in model.parameters() if p.requires_grad],
lr=config['training']['learning_rate'])
criterion = nn.CrossEntropyLoss(ignore_index=-100)
scheduler = torch.optim.lr_scheduler.CosineAnnealingLR(optimizer,
T_max=len(train_loader)*config['training']['num_epochs'])

print("\n--- Starting Training ---")
for epoch in range(config['training']['num_epochs']):
    train_loss = train_epoch(model, train_loader, optimizer, criterion, scheduler, 1.0,
config['model']['vocab_size'])
    val_loss = evaluate_epoch(model, val_loader, criterion, config['model']['vocab_size'], 1.0)

    # Log Gini
    if hasattr(model, "_orig_mod"):
        mid_layer =
model._orig_mod.layers[config['model']['num_layers'] // 2]['some']
    else:
        mid_layer = model.layers[config['model']['num_layers'] // 2]['some']

    gini = calculate_gini(mid_layer.usage_count)
    print(f"Epoch {epoch+1}: Train={train_loss:.4f}, Val={val_loss:.4f},
PPL={math.exp(val_loss):.2f}")
    print(f" Middle Layer Gini={gini:.3f}, Dead Experts={mid_layer.dead_expert_count.item()}")

if __name__ == "__main__":
    main()

```